

LAB2 | 2023

Matrix Multiplication (Multi-Threading)

مروان احمد ربيع احمد 19016605

1 Introduction

You are required to implement a multi-threaded matrix multiplication program.

The input to the program is two matrices A(xy) and B(yz) that are read from corresponding text files. The output is a matrix C(x*z) that is written to an output text file.

A parallelized version of matrix multiplication can be done using one of these three methods:

1. A thread computes the output C matrix i.e. without multi-threading. (A thread per matrix).

$$\begin{bmatrix} a_1 & a_2 & a_3 \\ a_4 & a_5 & a_6 \\ a_7 & a_8 & a_9 \end{bmatrix} \begin{bmatrix} b_1 & b_2 & b_3 \\ b_4 & b_5 & b_6 \\ b_7 & b_8 & b_9 \end{bmatrix} = \begin{bmatrix} c_1 & c_2 & c_3 \\ c_4 & c_5 & c_6 \\ c_7 & c_8 & c_9 \end{bmatrix}$$

Figure 1: Thread Per Matrix

2. A thread computes each row in the output C matrix. (A thread per row).

$$\begin{bmatrix} a_1 & a_2 & a_3 \\ a_4 & a_5 & a_6 \\ a_7 & a_8 & a_9 \end{bmatrix} \begin{bmatrix} b_1 & b_2 & b_3 \\ b_4 & b_5 & b_6 \\ b_7 & b_8 & b_9 \end{bmatrix} = \begin{bmatrix} c_1 & c_2 & c_3 \\ c_4 & c_5 & c_6 \\ c_7 & c_8 & c_9 \end{bmatrix}$$

Figure 2: Thread Per Row

3. A thread computes each element in the output C matrix. (A thread per element).

$$\begin{bmatrix} a_1 & a_2 & a_3 \\ a_4 & a_5 & a_6 \\ a_7 & a_8 & a_9 \end{bmatrix} \begin{bmatrix} b_1 & b_2 & b_3 \\ b_4 & b_5 & b_6 \\ b_7 & b_8 & b_9 \end{bmatrix} = \begin{bmatrix} c_1 & c_2 & c_3 \\ c_4 & c_5 & c_6 \\ c_7 & c_8 & c_9 \end{bmatrix}$$

Figure 3: Thread Per Element

2 Test Cases

2.1 Test Case 1

Method: A thread per element									
row=10 col=10									
415	430	445	460	475	490	505	520	535	550
940	980	1020	1060	1100	1140	1180	1220	1260	1300
1465	1530	1595	1660	1725	1790	1855	1920	1985	2050
1990	2080	2170	2260	2350	2440	2530	2620	2710	2800
2515	2630	2745	2860	2975	3090	3205	3320	3435	3550
3040	3180	3320	3460	3600	3740	3880	4020	4160	4300
3565	3730	3895	4060	4225	4390	4555	4720	4885	5050
4090	4280	4470	4660	4850	5040	5230	5420	5610	5800
4615	4830	5045	5260	5475	5690	5905	6120	6335	6550
5140	5380	5620	5860	6100	6340	6580	6820	7060	7300

Figure 4: Output of Test Case 1

```
Microseconds taken by method 1: 144
Microseconds taken by method 2: 1192
Microseconds taken by method 3: 4629
```

Figure 5: Time of Test Case 1

2.2 Test Case 2

```
Method: A thread per element
row=3 col=4
-1 10 -15 -28
-3 -10 15 -36
5 -2 -9 -20
```

Figure 6: Output of Test Case 2

```
Microseconds taken by method 1: 63
Microseconds taken by method 2: 773
Microseconds taken by method 3: 1265

```

Figure 7: Time of Test Case 2

2.3 Test Case 3

```
Method: A thread per element
row=5 col=4
175 190 205 220
400 440 480 520
625 690 755 820
850 940 1030 1120
1075 1190 1305 1420
```

Figure 8: Output of Test Case 3

```

Microseconds taken by method 1: 90
Microseconds taken by method 2: 580
Microseconds taken by method 3: 1389

```

Figure 9: Time of Test Case 3

3 Code

3.1 Includes and type definitions

```

1 #include <string.h>
2 #include <stdio.h>
3 #include <stdlib.h>
4 #include <pthread.h>
5 #include <sys/time.h>
6 // #define MAX_STR_SIZE 100
7 // int *lbl_matrix_1;
8 // int *lbl_matrix_2;
9 int *lbl_matrix_3_r_th;
10 int *lbl_matrix_3_e_th;
11
12 // int lbl_row_1 ,lbl_row_2 ,lbl_column_1 ,lbl_column_2 ;
13 char *lbl_output_filename ;
14
15
16
17 typedef struct
18 {
19     int row_1 ,row_2 ,col_1 ,col_2 ;
20     int* matrix_1 , *matrix_2 ;
21     int* matrix_3_tpm , *matrix_3_tpr , *matrix_3_tpe ;
22     int current_row ,current_column ;
23     struct timeval stop_tpm , start_tpm ;
24     struct timeval stop_tpr , start_tpr ;
25     struct timeval stop_tpe , start_tpe ;
26 }combined_matrix_struct ;
27 combined_matrix_struct lbl_matrices ;

```

3.2 Read the size of matrix

```

28
29 void read_matrix_size (int *row_par , int *column_par , char *file_name){
30     char* extension = ".txt";
31     char fileSpec [strlen(file_name)+strlen(extension)+1];
32     FILE *fp ;
33
34     sprintf( fileSpec , sizeof( fileSpec ) , "%s%s" , file_name ,
35             extension );
36
37     fp = fopen( fileSpec , "r" );
38     fscanf(fp , "row=%d col=%d" , row_par , column_par );
39     // printf("%d %d" , *row_par , *column_par )
40     fclose(fp );
41 }
```

3.3 Read the matrices from file

```
41 // void populate_matrix ( int *row_par , int *column_par , char *file_name
42 // , int matrix[*row_par][*column_par]) {
43 void populate_matrix ( int row_par , int column_par , char *file_name , int
44 *matrix) {
45     char* extension = ".txt";
46     char fileSpec [strlen(file_name)+strlen(extension)+1];
47     FILE *fp ;
48     // printf(" **\n%d, %d\n**\n" , row_par , column_par);
49     snprintf( fileSpec , sizeof( fileSpec ) , "%s%s" , file_name ,
50 extension );
51     fp = fopen( fileSpec , "r" );
52     fscanf(fp , "%*[^\n]\n");
53
54     for( int i=0;i<row_par ; i++)
55     {
56         for( int j=0;j<column_par ; j++)
57         {
58             // printf("%d, %d\n" , i ,j );
59             // int temp;
60             fscanf(fp , "%d " ,( matrix+i*column_par+j));
61             // printf("%d " ,temp );
62         }
63     }
64 }
```

3.4 Thread per matrix

```
63
64 void write_matrix_to_file(int *matrix ,int row ,int column , char *
65 filename , char* extension , char* message){
66     char fileSpec [strlen(filename)+strlen(extension)+1];
67     FILE *fp ;
68     snprintf( fileSpec , sizeof( fileSpec ) , "%s%s" , filename , extension
69 );
70     if(( fp=fopen( fileSpec , "w" ))==NULL) {
71         printf("Cannot open file .\n");
72     }
73     fwrite( message , sizeof(char) , strlen(message) , fp );
74     fprintf(fp , "\n");
75     fprintf(fp , "row=%d col=%d" , row ,column );
76     fprintf(fp , "\n");
77     for( int i=0; i<row ; i++){
78         for( int j=0; j<column ; j++){
79             fprintf(fp , "%d \t" , *(matrix + i * column+ j));
80         }
81         fprintf(fp , "\n");
82     }
83     fclose(fp );
84 }
85
86 void thread_per_matrix(combined_matrix_struct* lcl_matrices , char *
87 file_name){
88     gettimeofday(&glbl_matrices.start_tpm , NULL); //start checking time
89     // int matrix_3[*row_par_1][*column_par_2];
90     int out_row = lcl_matrices->row_1 ;
91     int out_col = lcl_matrices->col_2 ;
92     int inner_dim = lcl_matrices->col_1 ;
93     int *lcl_matrix_1 = lcl_matrices->matrix_1 ;
```

```

91 int *lcl_matrix_2 = lcl_matrices->matrix_2;
92 int *matrix_3 = lcl_matrices->matrix_3_tpm;
93
94 for (int i = 0; i < out_row; i++) {
95     for (int j = 0; j < out_col; j++) {
96         *(matrix_3 + i * out_col+ j) = 0;
97         for (int k = 0; k < inner_dim; k++) {
98             *(matrix_3 + i * out_col+ j) += *(lcl_matrix_1+i*
99             inner_dim+k) * *(lcl_matrix_2+k*out_col+j);
100        }
101    }
102    printf("\nThe matrix using thread per matrix is : \n");
103    for (int i = 0; i < out_row; i++) {
104        printf("\n");
105        for (int j = 0; j < out_col; j++) {
106            printf("%d\t",*(matrix_3 + i * out_col+ j));
107        }
108    }
109    printf("\n\n");
110    gettimeofday(&lbl_matrices.stop_tpm, NULL); //end checking time
111}

```

3.5 Thread per row

```

112 //Each thread computes single element in the resultant matrix
113 void *multiply_row(void* arg){
114     combined_matrix_struct *lcl_matrices;
115     lcl_matrices = (combined_matrix_struct *) arg;
116     // int current_row = (int *)arg;
117     int current_row = lcl_matrices->current_row;
118     int out_row = lcl_matrices->row_1;
119     int out_col = lcl_matrices->col_2;
120     int inner_dim = lcl_matrices->col_1;
121     int *lcl_matrix_1 = lcl_matrices->matrix_1;
122     int *lcl_matrix_2 = lcl_matrices->matrix_2;
123     int *matrix_3 = lbl_matrices.matrix_3_tpr;
124     // printf("%d ", out_col);
125     // printf("%d %d %d %d\n", out_row, out_col, inner_dim, current_row);
126     for (int j = 0; j < out_col; j++) {
127         // printf("%d",j);
128         *(matrix_3+current_row*out_col+j) = 0;
129         for (int k = 0; k < inner_dim; k++) {
130             *(matrix_3+current_row*out_col+j) += *(lcl_matrix_1+
131             current_row*inner_dim+k) * *(lcl_matrix_2+k*out_col+j);
132         }
133         // printf("%d", *(lbl_matrix_3_r_th+current_row*out_col+j));
134     }
135     pthread_exit(1);
136 }
137
138 void thread_per_row(void){
139     gettimeofday(&lbl_matrices.start_tpr, NULL); //start checking time
140     pthread_t threads[lbl_matrices.row_1];
141
142     for (int i = 0; i < lbl_matrices.row_1; i++) {
143         combined_matrix_struct *lcl_matrices;
144         lcl_matrices=(combined_matrix_struct *) malloc(sizeof(
145         combined_matrix_struct));
146         memcpy(lcl_matrices , &lbl_matrices , sizeof(

```

```

146     combined_matrix_struct));
147     lcl_matrices->current_row = i;
148     pthread_create(&threads[i], NULL, multiply_row, (void*)(lcl_matrices));
149 }
150
151 // joining and waiting for all threads to complete
152 for (int i = 0; i < glbl_matrices.row_1; i++)
153     pthread_join(threads[i], NULL);
154 printf("\nThe matrix using thrad per row is : \n");
155 for(int i=0;i<glbl_matrices.row_1;i++)
156 {
157     printf("\n");
158     for (int j=0;j<glbl_matrices.col_2;j++)
159     {
160         printf("%d\t",*(glbl_matrices.matrix_3_tpr + i *
161             glbl_matrices.col_2 + j));
162     }
163     printf("\n\n");
164     gettimeofday(&glbl_matrices.stop_tpr, NULL); //end checking time
165 }
```

3.6 Thread per element

```

166 void *multiply_element(void* arg){
167     combined_matrix_struct *lcl_matrices;
168     lcl_matrices = (combined_matrix_struct *) arg;
169     int current_row = lcl_matrices->current_row;
170     int current_column = lcl_matrices->current_column;
171     int out_row = lcl_matrices->row_1;
172     int out_col = lcl_matrices->col_2;
173     int inner_dim = lcl_matrices->col_1;
174     int *lcl_matrix_1 = lcl_matrices->matrix_1;
175     int *lcl_matrix_2 = lcl_matrices->matrix_2;
176     int *matrix_3 = glbl_matrices.matrix_3_tpe;
177
178     *(matrix_3+current_row*out_col+current_column) = 0;
179
180     for (int k = 0; k < inner_dim; k++) {
181         *(matrix_3+current_row*out_col+current_column) += *(lcl_matrix_1+current_row*inner_dim+k) * *(lcl_matrix_2+k*out_col+current_column);
182     }
183     pthread_exit(1);
184 }
185
186 void thread_per_element(void){
187     gettimeofday(&glbl_matrices.start_tpe, NULL); //start checking time
188     pthread_t threads[glbl_matrices.row_1][glbl_matrices.col_2];
189
190     for (int i = 0; i < glbl_matrices.row_1; i++) {
191         for (int j = 0; j < glbl_matrices.col_2; j++) {
192             combined_matrix_struct *lcl_matrices;
193             lcl_matrices=(combined_matrix_struct *)malloc(sizeof(combined_matrix_struct));
194             memcpy(lcl_matrices, &glbl_matrices, sizeof(combined_matrix_struct));
195             lcl_matrices->current_row = i;
196             lcl_matrices->current_column = j;
```

```

198     pthread_create(&threads[i][j], NULL, multiply_element, (
199         void *) (lcl_matrices));
200     }
201   }
202 
203   // joining and waiting for all threads to complete
204   for (int i = 0; i < glbl_matrices.row_1; i++) {
205     for (int j = 0; j < glbl_matrices.col_2; j++) {
206       pthread_join(threads[i][j], NULL);
207     }
208     printf("\nThe matrix using thread per element is : \n");
209   for (int i=0;i<glbl_matrices.row_1;i++)
210   {
211     printf("\n");
212     for (int j=0;j<glbl_matrices.col_2;j++)
213     {
214       printf("%d\t",*(glbl_matrices.matrix_3_tpe + i *
215 glbl_matrices.col_2 + j));
216     }
217     printf("\n\n");
218   gettimeofday(&glbl_matrices.stop_tpe, NULL); //end checking time
219 }
220 void main (int argc , char **argv){
```

3.7 Main

```

220 void main (int argc , char **argv){
221   // printf("You have entered %d arguments:\n", argc);
222   char *input_filename_1;
223   char *input_filename_2;
224   char *out_filename;
225 
226   if (argc == 4){
227     // char *input_filename_1,*input_filename_2;
228     // char *out_filename;
229     input_filename_1 = (char*) malloc((strlen(argv[1]) + 1) *
230 sizeof(char));
231     strcpy(input_filename_1 , argv[1]);
232     input_filename_2 = (char*) malloc((strlen(argv[2]) + 1) *
233 sizeof(char));
234     strcpy(input_filename_2 , argv[2]);
235     out_filename = (char*) malloc((strlen(argv[3]) + 1) * sizeof(
236 char));
237     strcpy(out_filename , argv[3]);
238   }
239   else{
240     input_filename_1 = (char*) malloc((strlen("a") + 1) * sizeof(
241 char));
242     strcpy(input_filename_1 , "a");
243     input_filename_2 = (char*) malloc((strlen("b") + 1) * sizeof(
244 char));
245     strcpy(input_filename_2 , "b");
246     out_filename = (char*) malloc((strlen("c") + 1) * sizeof(char));
247     strcpy(out_filename , "c");
248   }
249 
250   // printf("%s\n", input_filename_1);
251   // printf("%s\n", input_filename_2);
```

```

247 // printf("%s\n", out_filename);
248 read_matrix_size (&lbl_matrices.row_1, &lbl_matrices.col_1,
249 input_filename_1);
250 read_matrix_size (&lbl_matrices.row_2, &lbl_matrices.col_2,
251 input_filename_2);

252 lbl_matrices.matrix_1 = (int*) malloc(lbl_matrices.row_1 *
253 lbl_matrices.col_1 * sizeof(int));
254 lbl_matrices.matrix_2 = (int*) malloc(lbl_matrices.row_2 *
255 lbl_matrices.col_2 * sizeof(int));

256 populate_matrix (lbl_matrices.row_1, lbl_matrices.col_1,
257 input_filename_1, lbl_matrices.matrix_1);
258 populate_matrix (lbl_matrices.row_2, lbl_matrices.col_2,
259 input_filename_2, lbl_matrices.matrix_2);

260 lbl_matrices.matrix_3_tpm = (int*) malloc(lbl_matrices.row_1 *
261 lbl_matrices.col_2 * sizeof(int));
262 lbl_matrices.matrix_3_tpr = (int*) malloc(lbl_matrices.row_1 *
263 lbl_matrices.col_2 * sizeof(int));
264 lbl_matrices.matrix_3_tpe = (int*) malloc(lbl_matrices.row_1 *
265 lbl_matrices.col_2 * sizeof(int));

266 thread_per_matrix(&lbl_matrices, "c");
267 thread_per_row();
268 thread_per_element();

269 // printf("Seconds taken by method 1%lu\n", stop.tv_sec - start.
270 tv_sec);
271 printf("Microseconds taken by method 1: %lu\n", lbl_matrices.
272 stop_tpm.tv_usec - lbl_matrices.start_tpm.tv_usec);
273 printf("Microseconds taken by method 2: %lu\n", lbl_matrices.
274 stop_tpr.tv_usec - lbl_matrices.start_tpr.tv_usec);
275 printf("Microseconds taken by method 3: %lu\n", lbl_matrices.
276 stop_tpe.tv_usec - lbl_matrices.start_tpr.tv_usec);

277 write_matrix_to_file(lbl_matrices.matrix_3_tpm, lbl_matrices.
278 row_1, lbl_matrices.col_2, out_filename, "_per_matrix.txt", "
279 Method: A thread per matrix");
280 write_matrix_to_file(lbl_matrices.matrix_3_tpr, lbl_matrices.
281 row_1, lbl_matrices.col_2, out_filename, "_per_row.txt", "Method:
282 A thread per row");
283 write_matrix_to_file(lbl_matrices.matrix_3_tpe, lbl_matrices.
284 row_1, lbl_matrices.col_2, out_filename, "_per_element.txt", "
285 Method: A thread per element");

286 free(lbl_matrices.matrix_1);
287 free(lbl_matrices.matrix_2);
288 free(lbl_matrices.matrix_3_tpm);
289 free(lbl_matrices.matrix_3_tpr);
290 free(lbl_matrices.matrix_3_tpe);
291 }

```